

Learning Unitary Operators with Help From $\mathfrak{u}(n)$

@_hylandSL

ETH zürich

Introduction

- Training deep neural networks can be difficult due to the problem of **vanishing/exploding gradients**
- In recurrent neural networks, repeated application of the state evolution operator can be responsible
- Using a **norm-preserving** operator can help!
 - focus on unitary/orthogonal matrices
- Gradient updates in general break unitarity: use **parametrisation** closed under additive updates
- Idea: use **Lie group-Lie algebra correspondence** to define parametrisation

Parametrisation with $\mathfrak{u}(n)$

- Unitary $n \times n$ matrices form a continuous group, $U(n)$ (a **Lie group**), closed under matrix multiplication
- The tangent space to a Lie group at the identity is its **Lie algebra**: a vector space closed under *addition*
 - $U(n)$'s Lie algebra is $\mathfrak{u}(n)$, the space of $n \times n$ skew-Hermitian matrices ($A^\dagger = -A$)
- The **exponential map** associates elements of the Lie group and Lie algebra:
 - for matrix groups like $U(n)$, the map is the matrix exponential (**expm**)
 - $L \in \mathfrak{u}(n) \Rightarrow \exp(L) \in U(n)$
 - as $U(n)$ is compact and connected, the map is **surjective** (not true for the orthogonal group)
- Given a basis $\{T_i\}$ of $\mathfrak{u}(n)$, $L \in \mathfrak{u}(n)$ is defined by its coefficients $\lambda_i \in \mathbb{R}$
- Thus U is parametrised by these n^2 coefficients:
$$U(\lambda_1, \dots, \lambda_{n^2}) = \exp\left(\sum_{i=1}^{n^2} \lambda_i T_i\right)$$
 Thus: we do gradient descent in the **Lie algebra**
- Choice of basis $\{T_i\}$ is arbitrary: we choose them to be **sparse** (1 or 2 non-zero elements):
 - n diagonal imaginary
 - $n(n-1)/2$ symmetric imaginary
 - $n(n-1)/2$ skew-symmetric real


Stephanie L. Hyland, Gunnar Rätsch ETH Zürich, Weill Cornell Medicine

Learning the Operator

- Discard the rest of RNN machinery (for *now*!)
- Simple **supervised learning** task:
 - sample random ground truth U , many random \mathbf{x}
 - $\mathbf{y} = U\mathbf{x} + \epsilon$, where $\epsilon \sim \mathcal{N}(0, \sigma^2)$
 - minimise mean $\|\hat{U}\mathbf{x} - \mathbf{y}\|^2$ over batches
- Approaches for learning U :
 1. **projection**: after gradient update, project U back to closest unitary matrix
 2. **composition**: U is a product of parametrisable unitary operators; (with $7n$ free parameters) from Arjovsky, Shah & Bengio (ICLR 2016)
 3. **lie algebra**: U is defined by the n^2 coefficients of an element of the Lie algebra (my method)

n	true	projection	arjovsky	lie algebra	rand
3	$6.004 \pm 0.005 \times 10^{-4}$	8 ± 1	$6.005 \pm 0.003 \times 10^{-4}$	$6.003 \pm 0.003 \times 10^{-4}$	12.5 ± 0.4
6	~ 0.001	15 ± 1	0.09 ± 0.01	0.03 ± 0.01	24 ± 1
8	~ 0.002	14 ± 1	1.17 ± 0.06	0.014 ± 0.006	31.6 ± 0.6
14	~ 0.003	24 ± 4	10.8 ± 0.3	0.07 ± 0.02	52 ± 1
20	~ 0.004	38 ± 3	29.0 ± 0.5	0.47 ± 0.03	81 ± 2

Table: Mean squared ℓ_2 distance between true \mathbf{y} and predicted \mathbf{y}



Additional experiments:

- **restricted**: restrict lie algebra to have $7n$ learnable parameters
- **basis**: change the basis used by lie algebra with random uniform matrix (e.g. between -5 and 5)

Further Questions

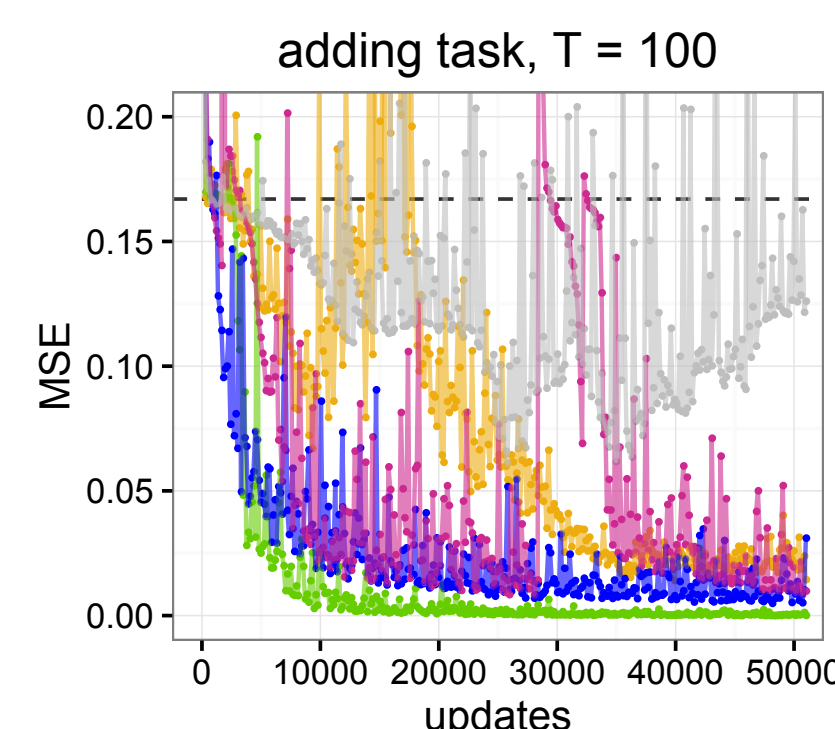
- **Geometry**:
 - Casting the learning problem directly as optimisation on the Lie group (manifold)
 - How does exponential map change cost surface?
- **Deep learning**:
 - understanding relative performance of models
 - nonlinearity has a large role in vanishing gradients - consider it **alongside** transition operator

General Unitary RNN

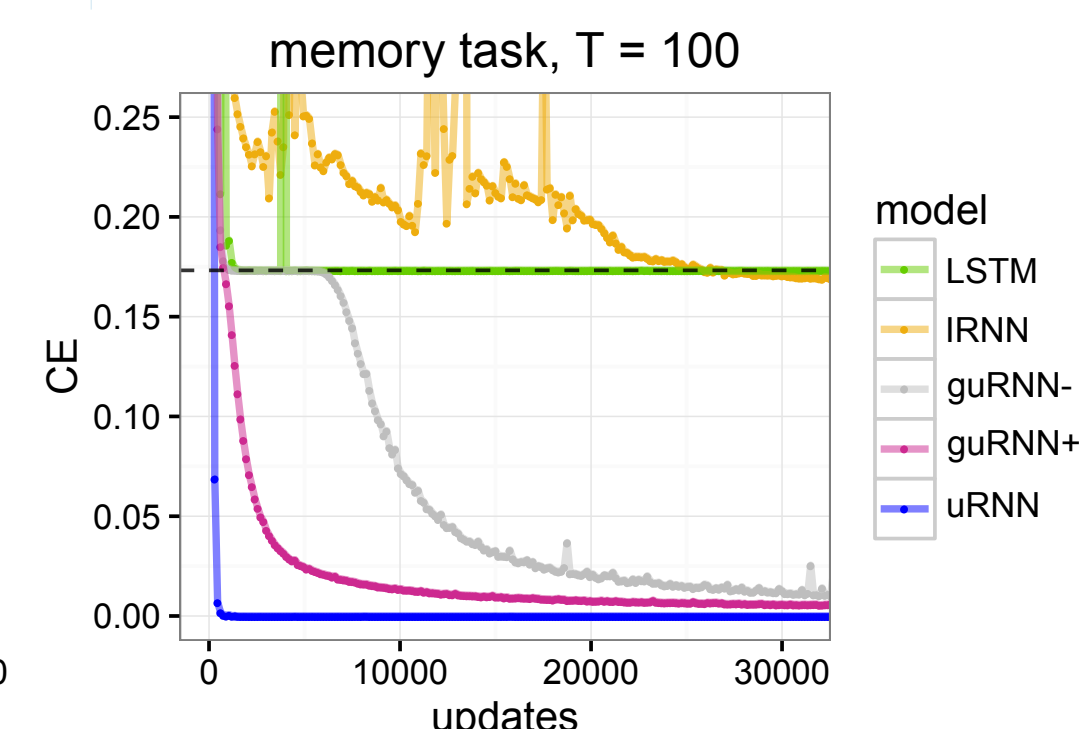
- Define (general) unitary recurrent neural network, with:
$$\mathbf{h}_t = f(U\mathbf{h}_{t-1} + V\mathbf{x}_t + \mathbf{b})$$
- f is a nonlinearity, \mathbf{h} is hidden state, U is a **unitary** matrix

Tasks: (classic long-memory RNN problems)

Adding problem: add two numbers from long sequence, minimise mean-squared error on result



Memory problem: 'remember' a sequence for a long time, minimise categorical cross-entropy over whole sequence



Conclusions

- Parametrisation using the **Lie algebra** enables the learning of arbitrary unitary operators with simple gradient descent
- We've used our parametrisation to define a **general unitary RNN** and demonstrated it on standard tasks
- The parametrisation of Arjovsky et al. struggles to learn matrices for $n > 7$, but still works well in an RNN setting, somehow

- **Code**: <https://github.com/corcra/uRNN>
- **Paper**: accepted at AAAI-17, available on arXiv:

